

# Lin Li

**Address:** Room 30.02, Wolfson Building, 14-15 Parks Rd, Oxford OX1 3QG, UK

**Email:** [lin.li@cs.ox.ac.uk](mailto:lin.li@cs.ox.ac.uk) | **Phone:** +86 189 0022 0595 (China) / +44 07536 139840 (UK)

**Web:** [treelli.github.io](https://treelli.github.io) | **Google Scholar:** [dxP6Y\\_oAAAAJ](https://scholar.google.com/citations?user=dxP6Y_oAAAAJ) | **Github:** [TreeLLi](https://github.com/TreELi)

## RESEARCH EXPERIENCE

---

**Postdoctoral Researcher**, Department of Computer Science, University of Oxford, UK 2025 – Present  
Advisor: [Prof. Yarin Gal](#). Group: [Oxford Applied and Theoretical Machine Learning \(OATML\)](#)

**Research Intern**, [Robotics X Lab](#), Tencent, Shenzhen, China 2021 – 22  
Mentor: [Prof. Lipeng Chen](#). Project: enabling robots to learn throwing and catching from human demonstrations.

## TEACHING EXPERIENCE

---

**Guest Lecturer**, Department of Computer Science, University of Oxford, Oxford, UK 2025  
Course: Uncertainty in Deep Learning (MSc, 20-30 students)

**Tutor**, Lady Margaret Hall, University of Oxford, Oxford, UK 2025  
Course (Summer School): Computer Vision (2 classes, 10-15 international students per class)

**Teaching Assistant**, Oxford and KCL 2021 / 25  
Courses: Uncertainty in Deep Learning (MSc), Machine Learning and Pattern Recognition (MSc), Introduction to Artificial Intelligence (Undergraduate)

## EDUCATION

---

**PhD in Machine Learning**, Department of Informatics, King's College London, UK 2019 – 24  
Advisor: [Prof. Michael Spratling](#). Thesis Committee: [Prof. George D. Magoulas](#) and [Dr. Adel Bibi](#).  
Thesis: *Towards Robust Visual Classification through Adversarial Training*

**MSc in Computing Science**, Department of Computing, Imperial College London, UK 2017 – 18  
Advisor: [Prof. Wayne Luk](#). Grade: Distinction (First Class Honour)

**BM in Financial Management**, School of Management, Xiamen University, China 2013 – 17  
Advisor: [Prof. Zheng Qiao](#). Grade: top 10% of graduating class

## PUBLICATION (\* indicates equal contribution; † indicates the corresponding author if have)

### AI Safety:

1. **Lin Li**†, Michael Spratling, "[Robust Shortcut and Disordered Robustness: Improving Adversarial Training Through Adaptive Smoothing](#)", the *Pattern Recognition (PR)*, 2025
2. **Lin Li**, Yifei Wang, Chawin Sitawarin, Michael Spratling, "[OODRobustBench: a Benchmark and Large-Scale Analysis of Adversarial Robustness under Distribution Shift](#)", the *International Conference on Machine Learning (ICML)*, 2024
3. **Lin Li**\*, Haoyan Guan\*, Jianing Qiu, Michael Spratling, "[One Prompt Word is Enough to Boost Adversarial Robustness for Pre-trained Vision-Language Models](#)", the *IEEE/CVF Computer Vision and Pattern Recognition Conference (CVPR)*, 2024
4. **Lin Li**†, Jianing Qiu, Michael Spratling, "[AROID: Improving Adversarial Robustness Through Online Instance-Wise Data Augmentation](#)", the *International Journal of Computer Vision (IJCV)*, 2024
5. **Lin Li**, Michael Spratling, "[Data augmentation alone can improve adversarial training](#)", the *International Conference on Learning Representations (ICLR)*, 2023

6. **Lin Li**, Michael Spratling†, "[Understanding and combating robust overfitting via input loss landscape analysis and regularization](#)", the *Pattern Recognition (PR)*, 2023
7. **Lin Li**, Georgia Channing, Suhaas M Bhat, Gabriel Davis Jones, Yarin Gal, "Building Reliable Long-Form Generation via Step-Wise Hallucination Rejection Sampling", *in submission*, 2025
8. Hengyang Yao\*, **Lin Li**\*, Ke Sun, Jianing Qiu, Huiping Chen, "Towards Robust Protective Perturbation against DeepFake Face Swapping", *in submission*, 2025
9. Jianing Qiu, **Lin Li**, Jiankai Sun, Hao Wei, Zhe Xu, Kyle Lam, Wu Yuan, "[Emerging Cyber Attack Risks of Medical AI Agents](#)", *in submission*, 2025

#### AI + X:

1. Jahan C Penny-Dimri, Magdalena Bachmann, William R Cooke, Sam Mathewlynn, Samuel Dockree, John Tolladay, Jannik Kossen, **Lin Li**, Yarin Gal, Gabriel Davis Jones, "[Measuring large language model uncertainty in women's health using semantic entropy and perplexity: a comparative study](#)", the *Lancet Obstetrics, Gynaecology, & Women's Health (Lancet OGW)*, 2025
2. Ravi Shankar, Sheng Wong, **Lin Li**, Magdalena Bachmann, Alex Silverthorne, Beth Albert, Gabriel Davis Jones, "[Energy Landscapes Enable Reliable Abstention in Retrieval-Augmented Large Language Models for Healthcare](#)", *NeurIPS Workshop GenAI4Health*, 2025
3. Lipeng Chen\*†, Jianing Qiu\*, **Lin Li**\*, Xi Luo, Guoyi Chi, Yu Zheng, "[Advancing Robots with Greater Dynamic Dexterity: A Large-Scale Multi-View and Multi-Modal Dataset of Human-Human Throw&Catch of Arbitrary Objects](#)", the *International Journal of Robotics Research (IJRR)*, 2024
4. Jianing Qiu\*, Jian Wu\*, Hao Wei\*, and Peilun Shi, Minqing Zhang, Yunyun Sun, **Lin Li**, ... (+34 authors), Wu Yuan†, "[Development and Validation of a Multimodal Multitask Vision Foundation Model for Generalist Ophthalmic Artificial Intelligence](#)", the *New England Journal of Medicine Artificial Intelligence (NEJM AI)*, 2024
5. Jianing Qiu, **Lin Li**, Jiankai Sun, and Jiachuan Peng, Peilun Shi, Ruiyang Zhang, Yinzhaodong, Kyle Lam, Frank P.-W. Lo, Bo Xiao, Wu Yuan, Dong Xu†, Benny Lo†, "[Large AI Models in Health Informatics: Applications, Challenges, and the Future](#)", the *IEEE Journal of Biomedical and Health Informatics (JBHI)*, 2023
6. Linyuan Li, Anujit Saha, **Lin Li**, Boyuan Li, Mengxian He, Jianing Qiu, Wu Yuan, "[Artificial Intelligence for Biomedical Video Generation](#)", *in submission*, 2024

#### GRANTS

**Advancing Large AI Models: Integration of New Data Modalities and Expansion of Capabilities** 2025-29  
Amount: €29M. Funder: [Horizon Europe](#)  
Role: Proposal contributor, and Research project leader. PI: Prof. Yarin Gal.

**Technical AI Safety Research** 2025-27  
Amount: \$371K. Funder: [Open Philanthropy](#)  
Role: Co-Principal Investigator (lead the grant application and the project). PI: Prof. Yarin Gal.

**Proactive Defense against Deepfake** 2025  
Amount: £20K. Funder: IDAI Pump Prime, University of Birmingham  
Role: Co-Investigator. PI: Prof. Huiping Chen.

#### AWARDS & HONOURS

|                                                 |         |
|-------------------------------------------------|---------|
| OpenAI API Access (Credits \$1000)              | 2025    |
| Nomination of USERN Prize                       | 2025    |
| KCL King's-China Scholarship (£109.8K)          | 2019-24 |
| KCL PGR Research Support (£2.5K)                | 2023    |
| Xiamen University First Class Scholarship (¥7K) | 2016    |

|                                                                    |      |
|--------------------------------------------------------------------|------|
| 3rd prize winner in Jinyuan Creativity and Startup Contest         | 2016 |
| Xiamen University Excellent Academic Performance Scholarship       | 2015 |
| 3rd prize winner in ChinaNet Dream Accelerator Programming Contest | 2015 |

## PRESS

- |    |                                                                                                                                                                                                       |      |
|----|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|------|
| 1. | <b>Süddeutsche Zeitung (KZ)</b> , major national newspapers, German<br>Article: <a href="#">Are Language Models Becoming Even More Unreliable?</a>                                                    | 2025 |
| 2. | <b>South China Morning Post (SCMP)</b> , major national newspapers, Hong Kong, China<br>Article: <a href="#">DeepSeek's AI in hospitals is 'too fast, too soon', Chinese medical researchers warn</a> | 2025 |
| 3. | <b>South China Morning Post (SCMP)</b> , major national newspapers, Hong Kong, China<br>Article: <a href="#">A Hong Kong AI model proves more accurate than doctors in diagnosing eye conditions</a>  | 2024 |

## TALKS

- |    |                                                                                                                                                                                                   |      |
|----|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|------|
| 1. | <b>AI Masterpiece Talk</b> , Outstanding Young Scholars Society, UK<br>Talk: <a href="#">Towards Robust and Safe AI: An Adversarial Approach</a> .<br>Host: Prof. Yu Guan and Prof. Sai Gu, FREng | 2025 |
| 2. | <b>AI Agents for Surgery Workshop</b> , 17th Hamlyn Symposium on Medical Robotics, UK<br>Talk: <a href="#">Hallucination of AI Agents</a> .                                                       | 2025 |
| 3. | <b>CHI Seminar</b> , Computational Health Informatics (CHI) Lab, University of Oxford, UK<br>Talk: Prompting VLMs for accuracy and robustness.<br>Host: Xiao Gu                                   | 2025 |
| 4. | <b>Nanqiang Youth Scholar Forum</b> , Xiamen University, China<br>Talk: <a href="#">Towards Safe Multimodal Foundational Models in Finance</a> .                                                  | 2024 |
| 5. | <b>AI Time Youth PhD Talk</b> , Tsinghua University, China.<br>Talk: Data augmentation can improve adversarial training.                                                                          | 2023 |
| 6. | <b>Departmental Research Showcase</b> , KCL, UK<br>Talk: Defending DNNs against adversarial examples.                                                                                             | 2023 |

## ACADEMIC SERVICE

### Reviewer

Conferences: NeurIPS, ICML, ICLR, CVPR, AAAI, ECCV.

Journals: IEEE T-Pattern Analysis and Machine Intelligence (T-PAMI), Pattern Recognition (PR), IEEE T-Information Forensics & Security (T-IFS), IEEE T-Dependable and Secure Computing (T-DSC), IEEE Journal of Biomedical and Health Informatics (J-BHI).

### Area Chair

Workshops: [Agentic AI for Medicine Workshop@MICCAI2025](#)

### Program Committee

Workshops: [2nd MEIS Workshop@CVPR2025](#), [WIHR@ICRA 2024](#), [Workshop on Cooperative Intelligence for Embodied AI@ECCV2024](#), [Agentic AI for Medicine Workshop@MICCAI2025](#)

## ACADEMIC REFEREES (refer to their websites for the latest information)

### Yarin Gal

Professor of Machine Learning, Department of Computer Science, University of Oxford.

Turing AI Fellow, The Alan Turing Institute, London, UK.

Expert Advisor, UK AI Security Institute (AISi), London, UK.

Tutorial Fellow, Christ Church College, University of Oxford.  
Phone: +44 1865 610497, Email: [yarin@cs.ox.ac.uk](mailto:yarin@cs.ox.ac.uk).  
Web: [www.cs.ox.ac.uk/people/yarin.gal/website](http://www.cs.ox.ac.uk/people/yarin.gal/website)

**Michael Spratling**

Researcher, Department of Behavioural and Cognitive Sciences, University of Luxembourg.  
was Reader<sup>1</sup> in Computational Neuroscience and Visual Cognition, Department of Informatics, KCL.  
Phone: +352 46 66 44 5722, Email: [michael.spratling@uni.lu](mailto:michael.spratling@uni.lu).  
Web: [mwspratling.codeberg.page](http://mwspratling.codeberg.page)

**Adel Bibi**

Senior Researcher (G9: equ. spinal point to Associate Professor) in ML, Department of Engineering Science, University of Oxford.  
Research Member of the Common Room at Kellogg College, University of Oxford.  
Phone: +44 7306 074827, Email: [adel.bibi@eng.ox.ac.uk](mailto:adel.bibi@eng.ox.ac.uk).  
Web: [www.adelbibi.com](http://www.adelbibi.com)

---

<sup>1</sup> Michael held this position while serving as Lin's PhD supervisor.